

Viraj Prabhu

Webpage: <https://virajprabhu.github.io>

Email: virajp@gatech.edu

RESEARCH INTERESTS

Domain Adaptation, Learning with limited supervision (Self-supervised, Active, & Few-shot Learning)

EDUCATION

Georgia Institute of Technology, Atlanta 2019 - present

Ph.D. in Computer Science, advised by Prof. Judy Hoffman

Georgia Institute of Technology, Atlanta 2017 - 2019

Master of Science in Computer Science, advised by Prof. Devi Parikh

GPA: 4.0, Awarded M.S. Research Award

Birla Institute of Technology and Science, Pilani 2011 - 2015

Bachelor of Engineering (with honors) in Computer Science

RESEARCH EXPERIENCE

Adaptive Learning Lab, Georgia Tech Fall 2019 - Present

Graduate Research Assistant, advised by Prof. Judy Hoffman Atlanta, GA

Developing data-efficient and reliable computer vision systems that can be deployed in the real world.

Salesforce Research Summer 2021

Research Intern, mentored by Dr. Nikhil Naik Palo Alto, CA

– Worked on adapting visual classifiers to unseen geographies.

Curai Summer 2018, 2019

Research Intern, mentored by Dr. Anitha Kannan Palo Alto, CA

– Developed open-set machine learning algorithms for disease diagnosis from clinical case data. Presented at ML4H at NeurIPS '19.

– Developed few-shot learning approach for dermatological diagnosis. Published at MLHC '19.

Visual Intelligence Lab, Georgia Tech Fall 2017 - Spring 2019

Graduate Research Assistant, advised by Prof. Devi Parikh Atlanta, GA

Worked on human-in-the-loop evaluation of visual conversational agents, and of “interpretability” modalities proposed for such agents. Published work at HCOMP '17 and EMNLP '18.

Machine Learning and Perception Lab, Virginia Tech Fall 2016 - Spring 2017

Research Assistant, advised by Prof. Dhruv Batra Blacksburg, VA

Worked on equipping VQA models with mechanisms for detecting the relevance of questions, and with better compositional reasoning. Published at EMNLP '17.

Adobe Summer 2014

Research Intern, Adobe Presenter Video Express (PVX) Bangalore, KA

Designed and implemented fast graphcut-based segmentation algorithm for real-time background substitution in video. Transferred into *Magic Green Screen*, the marquee feature of PVX 11.

AWARDS & SERVICE

Outstanding reviewer, NeurIPS 2021.

Outstanding reviewer, CVPR 2021.

M.S. Research Award, awarded by Georgia Tech's College of Computing (1 student annually).

Among *Top-30%* reviewers, NeurIPS 2018.

Reviewer, NeurIPS '18-'21, CVPR '18, '21-'22, ICLR '18,'20, ECCV '18, ACL '19.

1st, VTHacks 2017, Virginia Tech's annual hackathon. [Project].

1st, Google Hackathon, APOGEE 2014, for Snapify, an image-sharing app (from > 25 teams).

Subfinalist, LDV Entrepreneurial Computer Vision Challenge 2017, representing CloudCV.

2nd, Project Presentation, APOGEE 2013, for Try-On, a Kinect-based virtual dressing room app.

Awarded *Travel Scholarship*, for Google Summer of Code Mentor summit 2016, 2017.

Top-200 rank, BITSAT 2011, (from >120k applicants).

Awarded *Amul Vidya Shree* for *Top-100 rank* in ICSE 2009 (from >150k applicants).

PUBLICATIONS

Preprints

[13] **AUGCO: Augmentation Consistency-guided Self-training for Source-free Domain Adaptive Semantic Segmentation** [Paper]

V. Prabhu*, S. Khare*, D. Kartik, J. Hoffman. (2021) (* = equal)

Conference Papers

[12] **Mitigating Bias in Visual Transformers via Targeted Alignment**

S. Sudhakar, **V. Prabhu**, A. Krishnakumar, J. Hoffman. (2021) *British Machine Vision Conference (BMVC) 2021*

[11] **UDIS: Unsupervised Discovery of Bias in Deep Visual Recognition Models**

A. Krishnakumar, **V. Prabhu**, S. Sudhakar, J. Hoffman. (2021) *British Machine Vision Conference (BMVC) 2021*.

[10] **SENTRY: Selective Entropy Optimization via Committee Consistency for Unsupervised Domain Adaptation** [Project Page]

V. Prabhu, S. Khare, D. Kartik, J. Hoffman.

International Conference on Computer Vision (ICCV) 2021.

[9] **Active Domain Adaptation via Clustering Uncertainty-weighted Embeddings** [Project Page]

V. Prabhu, A. Chandrasekaran, K. Saenko, J. Hoffman.

International Conference on Computer Vision (ICCV) 2021.

[8] **Few-Shot Learning for Dermatological Disease Diagnosis.** [Paper][Poster]

V. Prabhu, A. Kannan, M. Ravuri, M. Chablani, D. Sontag, X. Amatriain.

Machine Learning and Healthcare Conference, 2019 (Spotlight).

[7] **Do Explanations make VQA Models more Predictable to a Human?** [Paper]

A. Chandrasekaran*, **V. Prabhu***, D. Yadav*, P. Chattopadhyay*, D. Parikh.

Conference on Empirical Methods in Natural Language Processing (EMNLP) 2018. (* = equal)

[6] **The Promise of Premise: Harnessing Question Premises in Visual Question Answering.** [Paper]

A. Mahendru*, **V. Prabhu***, A. Mohapatra*, D. Batra, S. Lee.

Conference on Empirical Methods in Natural Language Processing (EMNLP) 2017.

[5] **Evaluating Visual Conversational Agents via Cooperative Human-AI Games.** [Paper]

P. Chattopadhyay*, D. Yadav*, **V. Prabhu**, A. Chandrasekaran, A. Das, S. Lee, D. Batra, D. Parikh.

AAAI Conference on Human Computation and Crowdsourcing (HCOMP) 2017.

Workshop Papers

[4] **Open Set Medical Diagnosis** [Paper]

V. Prabhu, A. Kannan, G. Tso, N. Katariya, M. Chablani, D. Sontag, X. Amatriain.
ML for Health Workshop, NeurIPS 2019.

[3] **Fabrik: An Online Collaborative Neural Network Editor.** [Paper]

U. Garg, **V. Prabhu**, D. Yadav, R. Ramrakhya, H. Agarwal, D. Batra.
Workshop on AI Systems, SOSIP 2019.

[2] **Few-Shot Learning for Dermatological Disease Diagnosis.**

V. Prabhu, A. Kannan, M. Ravuri, M. Chablani, D. Sontag, X. Amatriain.
ML for Health Workshop, NeurIPS 2018.

[1] **It Takes Two to Tango: Towards Theory of AI's Mind.** [Paper]

A. Chandrasekaran*, D. Yadav*, P. Chattopadhyay*, **V. Prabhu***, D. Parikh.
Chalearn Looking at People Workshop, CVPR 2017 (Oral).

PROGRAMMING EXPERIENCE

CloudCV

Summer 2016, 2017

Mentor, Google Summer of Code, Google Code-In

Blacksburg, VA

Lead mentor for Fabrik (github.com/Cloud-CV/Fabrik), an open-source web platform to collaboratively build, visualize, and design neural networks in the browser. (**1000+** **GitHub stars**)

Adobe Systems

2015 - 2016

Member of Technical Staff, Adobe Captivate Prime

Bangalore, KA

Owner of Captivate Prime Android app for two release cycles. Developed features for offline playback, sync, and internationalization.

Tonbo Imaging

Fall 2016 - Spring 2017

R&D Intern

Bangalore, KA

Automated Calibration: Developed algorithm for automated calibration of company cameras using a collimator and AprilTag target setup, reducing calibration error by 6%.

Boresighting: Developed a boresighting algorithm to precisely align a weapon's muzzle and sighting system with a target at 10m to 100m for TDS-BRS, Tonbo's video precision boresight tool.

CEERI Pilani

Spring 2014

Project Assistant, advised by Prof. Jagdish Raheja

Pilani, RA

Built Kinect-based teleconferencing app that detected and displayed the current speaker.

TEACHING EXPERIENCE

Head Teaching Assistant, Intro to Computer Vision

Spring 2021

Course Instructor: Prof. Judy Hoffman

Georgia Tech

Worked with instructor and team of 5 TA's to conduct the course. Also responsible for designing homeworks, weekly office hours, and grading.

Teaching Assistant, Deep Learning

Fall 2019

Course Instructor: Prof. Dhruv Batra

Georgia Tech

Taught lecture on Reinforcement Learning ([Slides](#)). Held weekly hours, and graded homeworks.

Teaching Assistant, Intro to Machine Learning

Fall 2016

Course Instructor: Prof. Stefan Lee

Virginia Tech

Created homework machine learning challenges on Kaggle, and graded homeworks.

OTHER PROJECTS

Learning Cooperative Visual Dialog Agents via Deep Reinforcement Learning

PyTorch implementation (**130+ GitHub stars**) of Das & Kottur et al, ICCV '17. Used as starting point for starter code for 2018 Visual Dialog challenge. (github.com/batra-mlp-lab/visdial-rl)

Learning Active Learning Policies for Visual Recognition

[Report]

Course Project, Adaptive Control and Reinforcement Learning

Spring 2019

Explored strategies to learn active learning policies for visual recognition via reinforcement learning.

Visual Dialog Models that Rollout a Mental Model of their Interlocutors

[Poster]

Course Project, Deep Learning

Fall 2017

Explored self-play strategies based on dialog rollouts to develop cooperative visual dialog agents.

Exploring Weak Supervision and Generative Models for Semantic Segmentation

[Report]

Course Project, Probabilistic Graphical Models

Spring 2018

- Explored weakly supervised semantic segmentation using localization cues from GradCAM.
- Studied semantic segmentation via deep probabilistic generative models.

SELECTED COURSEWORK

Graduate: Adaptive Control and Reinforcement Learning, Advanced Machine Learning, Probabilistic Graphical Models, Machine Learning, Deep Learning, Computer Vision, Computability & Algorithms, Information Visualization, High-dimensional Data Analytics

Undergraduate: Pattern Recognition, Information Retrieval, Parallel Computing, Operating Systems, Advanced Algorithms, Computer Architecture, Computer Networks

PROGRAMMING SKILLS

Languages: Python, Lua, C/C++, Java, JavaScript, MATLAB, Shell

Technologies: PyTorch, Keras, L^AT_EX, HTML/CSS, ReactJS, EmberJS, Android

REFERENCES

- Prof. Judy Hoffman, Georgia Tech (email: judy@gatech.edu)
- Prof. Devi Parikh, Georgia Tech (email: parikh@gatech.edu)
- Prof. Dhruv Batra, Georgia Tech (email: dbatra@gatech.edu)
- Dr. Anitha Kannan, Curai (email: anitha@curai.com)
- Prof. Stefan Lee, Oregon State University (email: leestef@oregonstate.edu)