Viraj Prabhu

Webpage: https://virajprabhu.github.io

Email: virajp@gatech.edu
Phone: *(470)-494-1837*

RESEARCH INTERESTS

Transfer learning, Learning with limited supervision, Uncertainty Estimation, Vision & Language

EDUCATION

Georgia Institute of Technology, Atlanta   *2019 - present*
Ph.D. in Computer Science, advised by Prof. Judy Hoffman

Georgia Institute of Technology, Atlanta   *2017 - 2019*
Master of Science in Computer Science, advised by Prof. Devi Parikh
GPA: 4.0, Awarded M.S. Research Award

Birla Institute of Technology and Science, Pilani   *2011 - 2015*
Bachelor of Engineering (with honors) in Computer Science

RESEARCH EXPERIENCE

Curai   Summer 2018, 2019
Research Intern, mentored by Dr. Anitha Kannan   *Palo Alto, CA*

− Developed open-set machine learning algorithms for disease diagnosis from clinical case data. To appear at ML4H at NeurIPS '19.
− Developed few-shot learning approach for dermatological diagnosis. Published at MLHC '19.

Visual Intelligence Lab, Georgia Tech   Fall 2017 - Spring 2019
Graduate Research Assistant, advised by Prof. Devi Parikh   *Atlanta, GA*

Worked on human-in-the-loop evaluation of visual conversational agents, and of "interpretability" modalities proposed for such agents. Published work at HCOMP '17 and EMNLP '18.

Machine Learning and Perception Lab, Virginia Tech   Fall 2016 - Spring 2017
Research Assistant, advised by Prof. Dhruv Batra   *Blacksburg, VA*

Worked on equipping VQA models with mechanisms for detecting the relevance of questions, and with better compositional reasoning. Published at EMNLP '17.

Adobe   Summer 2014
Research Intern, Adobe Presenter Video Express (PVX)   *Bangalore, KA*

Designed and implemented fast graphcut-based segmentation algorithm for real-time background substitution in video. Transferred into *Magic Green Screen*, the marquee feature of PVX 11.

AWARDS & SERVICE

M.S. Research Award, awarded by Georgia Tech's College of Computing (1 student annually).
Among *Top-30%* reviewers, NeurIPS 2018.
Reviewer, NeurIPS, ACL 2019, NeurIPS, CVPR, ICLR, ECCV 2018.
1st, VTHacks 2017, Virginia Tech's annual hackathon. [Project].
1st, Google Hackathon, APOGEE 2014, for Snapify, an image-sharing app (from > 25 teams).
Subfinalist, LDV Entrepreneurial Computer Vision Challenge 2017, representing CloudCV.
2nd, Project Presentation, APOGEE 2013, for Try-On, a Kinect-based virtual dressing room app.
Awarded *Travel Scholarship*, for Google Summer of Code Mentor summit 2016, 2017.
Top-200 rank, BITSAT 2011, (from >120k applicants).
Awarded *Amul Vidya Shree* for *Top-100 rank* in ICSE 2009 (from >150k applicants).

PUBLICATIONS

Conference Papers

[8] **Few-Shot Learning for Dermatological Disease Diagnosis.** [Paper][Poster]
V. Prabhu, A. Kannan, M. Ravuri, M. Chablani, D. Sontag, X. Amatriain.
*Machine Learning and Healthcare Conference, 2019 (**Spotlight**).*

[7] **Do Explanations make VQA Models more Predictable to a Human?** [Paper]
A. Chandrasekaran*, **V. Prabhu***, D. Yadav*, P. Chattopadhyay*, D. Parikh.
Conference on Empirical Methods in Natural Language Processing (EMNLP) 2018. (* = equal)

[6] **The Promise of Premise: Harnessing Question Premises in Visual Question Answering.** [Paper]
A. Mahendru*, **V. Prabhu***, A. Mohapatra*, D. Batra, S. Lee.
Conference on Empirical Methods in Natural Language Processing (EMNLP) 2017.

[5] **Evaluating Visual Conversational Agents via Cooperative Human-AI Games.** [Paper]
P. Chattopadhyay*, D. Yadav*, **V. Prabhu**, A. Chandrasekaran, A. Das, S. Lee, D. Batra, D. Parikh.
AAAI Conference on Human Computation and Crowdsourcing (HCOMP) 2017.

Workshop Papers

[4] **Open Set Medical Diagnosis** [Paper]
V. Prabhu, A. Kannan, G. Tso, N. Katariya, M. Chablani, D. Sontag, X. Amatriain.
ML for Health Workshop, NeurIPS 2019.

[3] **Fabrik: An Online Collaborative Neural Network Editor.** [Paper]
U. Garg, **V. Prabhu**, D. Yadav, R. Ramrakhya, H. Agarwal, D. Batra.
Workshop on AI Systems, SOSP 2019.

[2] **Few-Shot Learning for Dermatological Disease Diagnosis.**
V. Prabhu, A. Kannan, M. Ravuri, M. Chablani, D. Sontag, X. Amatriain.
ML for Health Workshop, NeurIPS 2018.

[1] **It Takes Two to Tango: Towards Theory of AI's Mind.** [Paper]
A. Chandrasekaran*, D. Yadav*, P. Chattopadhyay*, **V. Prabhu***, D. Parikh.
*Chalearn Looking at People Workshop, CVPR 2017 (**Oral**).*

PROGRAMMING EXPERIENCE

CloudCV   Summer 2016, 2017
Mentor, Google Summer of Code, Google Code-In   *Blacksburg, VA*

Lead mentor for Fabrik (`github.com/Cloud-CV/Fabrik`), an open-source web platform to collaboratively build, visualize, and design neural networks in the browser. (**1000+ GitHub stars**)

Adobe Systems   2015 - 2016
Member of Technical Staff, Adobe Captivate Prime   *Bangalore, KA*

Owner of Captivate Prime Android app for two release cycles. Developed features for offline playback, sync, and internationalization.

Tonbo Imaging   Fall 2016 - Spring 2017
R&D Intern   *Bangalore, KA*

Automated Calibration: Developed algorithm for automated calibration of company cameras using a collimator and AprilTag target setup, reducing calibration error by 6%.

Boresighting: Developed a boresighting algorithm to precisely align a weapon's muzzle and sighting system with a target at 10m to 100m for TDS-BRS, Tonbo's video precision boresight tool.

CEERI Pilani Spring 2014
Project Assistant, advised by Prof. Jagdish Raheja *Pilani, RA*

Built Kinect-based teleconferencing app that detected and displayed the current speaker.

TEACHING EXPERIENCE

Teaching Assistant, Deep Learning Fall 2019
Course Instructor: Prof. Dhruv Batra *Georgia Tech*

Taught lecture on Reinforcement Learning (Slides). Held weekly hours, and graded homeworks.

Teaching Assistant, Intro to Machine Learning Fall 2016
Course Instructor: Prof. Stefan Lee *Virginia Tech*

Created homework machine learning challenges on Kaggle, and graded homeworks.

OTHER PROJECTS

Learning Cooperative Visual Dialog Agents via Deep Reinforcement Learning

PyTorch implementation (**130+ GitHub stars**) of Das & Kottur et al, ICCV '17. Used as starting point for starter code for 2018 Visual Dialog challenge. (`github.com/batra-mlp-lab/visdial-rl`)

Learning Active Learning Policies for Visual Recognition [Report]
Course Project, Adaptive Control and Reinforcement Learning *Spring 2019*

Explored strategies to learn active learning policies for visual recognition via reinforcement learning.

Visual Dialog Models that Rollout a Mental Model of their Interlocutors [Poster]
Course Project, Deep Learning *Fall 2017*

Explored self-play strategies based on dialog rollouts to develop cooperative visual dialog agents.

Exploring Weak Supervision and Generative Models for Semantic Segmentation [Report]
Course Project, Probabilistic Graphical Models *Spring 2018*

– Explored weakly supervised semantic segmentation using localization cues from GradCAM.
– Studied semantic segmentation via deep probabilistic generative models.

SELECTED COURSEWORK

Graduate: Adaptive Control and Reinforcement Learning, Probabilistic Graphical Models, Machine Learning, Deep Learning, Computer Vision, Computability & Algorithms, Information Visualization
Undergraduate: Pattern Recognition, Information Retrieval, Parallel Computing, Operating Systems, Advanced Algorithms, Computer Architecture, Computer Networks

PROGRAMMING SKILLS

Languages: Python, Lua, C/C++, Java, JavaScript, MATLAB, Shell

Technologies: PyTorch, Torch, Keras, TensorFlow, LaTeX, ReactJS, EmberJS, Android

REFERENCES

- Prof. Judy Hoffman, Georgia Tech (email: judy@gatech.edu)
- Prof. Devi Parikh, Georgia Tech (email: parikh@gatech.edu)
- Prof. Dhruv Batra, Georgia Tech (email: dbatra@gatech.edu)
- Dr. Anitha Kannan, Curai (email: anitha@curai.com)
- Prof. Stefan Lee, Oregon State University (email: leestef@oregonstate.edu)